

AMATH 483 / 583 - HW3

1 Problems

1. Matrix Multiplication Loop Permutations. Implement templated gemm for each $\{i, j, k\}$ loop permutation using the following specifications. Each computes $C \leftarrow \alpha AB + \beta C$, where $A \in \mathbb{R}^{m \times p}$, $B \in \mathbb{R}^{p \times n}$, $C \in \mathbb{R}^{m \times n}$, $\alpha, \beta \in \mathbb{R}$, but will exhibit distinct memory access patterns. Check these produce the correct results. Turn in the .cpp and .hpp files for each. Include the header files into another header file `hw3_p1_header.hpp` and submit this as well. Pay special attention that your matrices will now be represented within a single vector in this exercise. Please utilize column major ordering as discussed in lecture when assigning and accessing matrix elements in this format.

- **template<typename T>**
void mm_ijk(**T** a, **const** std::vector<**T**>& A, **const** std::vector<**T**>& B, **T** b, std::vector<**T**>& C, **int** m, **int** p, **int** n);
- **template<typename T>**
void mm_jki(**T** a, **const** std::vector<**T**>& A, **const** std::vector<**T**>& B, **T** b, std::vector<**T**>& C, **int** m, **int** p, **int** n);
- **template<typename T>**
void mm_kij(**T** a, **const** std::vector<**T**>& A, **const** std::vector<**T**>& B, **T** b, std::vector<**T**>& C, **int** m, **int** p, **int** n);
- **template<typename T>**
void mm_jik(**T** a, **const** std::vector<**T**>& A, **const** std::vector<**T**>& B, **T** b, std::vector<**T**>& C, **int** m, **int** p, **int** n);
- **template<typename T>**
void mm_ikj(**T** a, **const** std::vector<**T**>& A, **const** std::vector<**T**>& B, **T** b, std::vector<**T**>& C, **int** m, **int** p, **int** n);
- **template<typename T>**
void mm_kji(**T** a, **const** std::vector<**T**>& A, **const** std::vector<**T**>& B, **T** b, std::vector<**T**>& C, **int** m, **int** p, **int** n);

2. Compiler Optimization. Use the $\{kij\}$ and $\{jki\}$ loop permutation codes from problem 1 to explore the performance of your implementations applying compiler optimization levels *default* (no optimization or default case), *-O3*, and *-ffast-math* (or the equivalent for your compiler!) for square matrices of dimension $n = 2$ to $n = 512$, stride one. Let each n be measured $ntrial$ times and plot the average performance for each case versus n , $ntrial \geq 3$. Submit your .cpp test code, and two plots -one for each loop variant on your choice of data type *float* or *double*. **Extra credit:** Submit plots for both data types.
3. (AM583 only, +5 for AM483) Strassen. Use notes from the class lecture to implement a C++ template for the (recursive) Strassen matrix multiplication algorithm. Plot the *double* precision performance for square matrices of even dimension from $n = 2$ to $n = 512$. Let each n be measured $ntrial$ times and plot the average performance versus n , $ntrial \geq 3$. You will turn in the .cpp and .hpp files for the strassen code, your .cpp test code, and performance plot.

```
template <typename T>
vector<vector<T>> strassen_mm(const vector<vector<T>> &A,
const vector<vector<T>> &B); //vector<vector<double>>C=strassen_mm(A, B);
```